

# Final Project Presentation MP2 Calculation

Team 10  
Hieu, Mouza, Olga, Artur





# Intro

Step 1: Change of basis  $(pq|P) = \sum_q^{N+M} (C^\dagger)_{qv} \sum_p^{N+M} (C^\dagger)_{p\mu} (\mu\nu|P)$

Step 2: Obtain B\_matrix  
(Orthogonalization)  $B_{pq}^P = \sum_P^{N+M} (pq|P) (P|Q)^{-1/2}$

$$(ia|jb) = \sum_P B_{ia}^P B_{jb}^P$$




Step 3: Calculate MP2

$$E_{MP2} = \sum_{ij} \sum_{ab} \frac{(ia | jb)[2(ia | jb) - (ib | ja)]}{\varepsilon_i + \varepsilon_j - \varepsilon_a - \varepsilon_b}$$

$(C^\dagger)$	NBASIS NBASIS
$(\mu\nu P)$	NBASIS NBASIS NAUX
$(P Q)^{-1/2}$	NAUX NAUX
MO_energy	NBASIS

# Systems of Interest

	$C_3H_8$	$C_{60}$	Heme
$N_{\text{BASIS}}$	202	1800	632
$N_{\text{AUX}}$	483	4860	1618
$N_{\text{ELEC}}$	13	180	161
Memory	157 MB	110 GB	5.17 GB
Total node hours	2.96 s	6.61 hours	0.0903 hours





# Serial Implementation

Use BLAS Level 3 DGEMM to optimize serial code

1)

```
MO1 and MO2
1 for (size_t P = 0; P < NAUX; P++){
2     cblas_dgemm(CblasRowMajor, CblasNoTrans, NBASIS, NBASIS, NBASIS, 1.0,
3         C, NBASIS, uqP + P * NBASIS * NBASIS, NBASIS, 1.0, pqP + P * NBASIS * NBASIS, NBASIS);
4 }
```

2)

3)

$$(pq|P) = \sum_q^{N+M} (C^\dagger)_{q\nu} \sum_p^{N+M} (C^\dagger)_{p\mu} (\mu\nu|P)$$



# Serial Implementation

Use BLAS Level 3 DGEMM to optimize serial code

1)

```
MO1 and MO2
1 for (size_t P = 0; P < NAUX; P++){
2     cblas_dgemm(CblasRowMajor, CblasNoTrans, CblasNoTrans, NBASIS, NBASIS, NBASIS, 1.0,
3         C, NBASIS, uqP + P * NBASIS * NBASIS, NBASIS, 1.0, pqP + P * NBASIS * NBASIS, NBASIS);
4 }
```

2)

```
Orthogonalization
1 cblas_dgemm(CblasColMajor, CblasNoTrans, CblasNoTrans, NBASIS*NBASIS, NAUX, NAUX, 1.0,
2     pqP, NBASIS*NBASIS, overlap, NAUX, 1.0, B_pqP, NBASIS*NBASIS);
```

3)

```
Orthogonalization
1 for (size_t p = 0; p < NBASIS; p++){
2     cblas_dgemm(CblasRowMajor, CblasNoTrans, CblasNoTrans, NBASIS, NAUX, NAUX, 1.0,
3         pqP_new+p*NAUX*NBASIS, NAUX, overlap, NAUX, 1.0, B_pqP+p*NAUX*NBASIS, NAUX);
4 }
```

$$B_{pq}^P = \sum_P^{N+M} (pq|P) (P|Q)^{-1/2}$$



# Serial Implementation

Use BLAS Level 3 DGEMM to optimize serial code

1)

```
MO1 and MO2

1 for (size_t P = 0; P < NAUX; P++){
2     cblas_dgemm(CblasRowMajor, CblasNoTrans, CblasNoTrans, NBASIS, NBASIS, NBASIS, 1.0,
3                 C, NBASIS, uqP + P * NBASIS * NBASIS, NBASIS, 1.0, pqP + P * NBASIS * NBASIS, NBASIS);
4 }
```

2)

```
Orthogonalization

1 cblas_dgemm(CblasColMajor, CblasNoTrans, CblasNoTrans, NBASIS*NBASIS, NAUX, NAUX, 1.0,
2             pqP, NBASIS*NBASIS, overlap, NAUX, 1.0, B_pqP, NBASIS*NBASIS);
```

```
Orthogonalization

1 for (size_t p = 0; p < NBASIS; p++){
2     cblas_dgemm(CblasRowMajor, CblasNoTrans, CblasNoTrans, NBASIS, NAUX, NAUX, 1.0,
3                 pqP_new + p * NAUX * NBASIS, NAUX, overlap, NAUX, 1.0, B_pqP + p * NAUX * NBASIS, NAUX);
4 }
```

3)

```
MP2 calculation

1 for (size_t i = 0; i < NELECS; i++){
2     for (size_t j = 0; j < NELECS; j++){
3         int NVIR = NBASIS - NELECS;
4         double *ipjq = new double[NVIR * NVIR] ();
5         // contract over P
6         cblas_dgemm(CblasRowMajor, CblasTrans, CblasNoTrans, NVIR, NVIR, NAUX, 1.0, B_iaP +
7                     i * NVIR * NAUX, NVIR, B_iaP + j * NVIR * NAUX, NVIR, 1.0, ipjq, NVIR);
8
9         // Vectorization can be done here since this is matrix matrix element wise multiplication
10        for (size_t a = NELECS; a < NBASIS; a++){
11            for (size_t b = NELECS; b < NBASIS; b++){
12                double ibja = ipjq[(b - NELECS) * NVIR + (a - NELECS)];
13                double iajb = ipjq[(a - NELECS) * NVIR + (b - NELECS)];
14                double delta = MO_energy[i] + MO_energy[j] - MO_energy[a] - MO_energy[b];
15                E_MP2 += 1/delta * (2 * iajb * iajb - iajb * ibja);
16            }
17        }
18    }
19 }
```

$$(ia|jb) = \sum_P B_{ia}^P B_{jb}^P$$


$$E_{MP2} = \sum_{ij} \sum_{ab} \frac{(ia|jb)[2(ia|jb) - (ib|ja)]}{\varepsilon_i + \varepsilon_j - \varepsilon_a - \varepsilon_b}$$



# Roofline Analysis

$$\pi = 1075.2 \text{ Gflop / s}$$

$$\beta = 76.8 \text{ GB / s}$$





# Roofline Analysis

$$F_{\text{MO1}, 2} = 2 \times N_{\text{BASIS}}^3 \times N_{\text{AUX}}$$


$$F_{\text{Orthog}} = 2 \times N_{\text{AUX}}^3 \times N_{\text{BASIS}}$$

$$F_{\text{MP2}} = 2 \times (N_{\text{BASIS}} - N_{\text{ELECS}})^2 \times N_{\text{AUX}} + (N_{\text{BASIS}} - N_{\text{ELECS}})^2 \times 11)$$

$$I_{\text{MO1}, 2} = \frac{F_{\text{MO1}, 2}}{8 \cdot 4 \cdot N_{\text{BASIS}}^2 N_{\text{AUX}}} = \frac{1}{16} N_{\text{BASIS}}$$

$$I_{\text{Orthog}} = \frac{F_{\text{Orthog}}}{8 \cdot 4 N_{\text{AUX}}^2 N_{\text{BASIS}}} = \frac{1}{16} N_{\text{AUX}}$$

$$I_{\text{MP2}} = \frac{F_{\text{MP2}}}{3 N_{\text{BASIS}}^2} = \frac{1}{16} N_{\text{AUX}}$$







# Roofline Analysis: BLAS routines


$$F_{\text{MO1}, 2} = 2 \times N_{\text{BASIS}}^3 \times N_{\text{AUX}}$$

$$F_{\text{Orthog}} = 2 \times N_{\text{AUX}}^3 \times N_{\text{BASIS}}$$

$$F_{\text{MP2}} = 2 \times (N_{\text{BASIS}} - N_{\text{ELECS}})^2 \times N_{\text{AUX}} + (N_{\text{BASIS}} - N_{\text{ELECS}})^2 \times 11)$$

\* cache optimization

Step	for-loops	CBLAS	OpenMP
MO1	20.99	1.42	0.09
MO2	28.67	1.45	0.13
Orthogonalization	50.54	3.29	0.28
MP2	34.04	1.18	0.07






# Performance Benchmarks: OpenMP implementation / strong scaling

#pragma omp parallel for

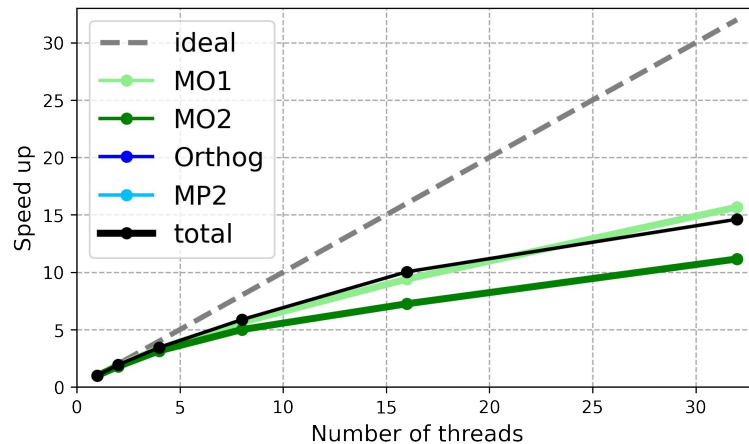

- scheduling
- **first touch policy**
- **thread affinity**



# Performance Benchmarks: OpenMP implementation / strong scaling

#pragma omp parallel for

- MO2 has worse cache locality





# Performance Benchmarks: OpenMP implementation / strong scaling

#pragma omp parallel for collapse(3)

- reshaping

#pragma omp parallel for

- DGEMM



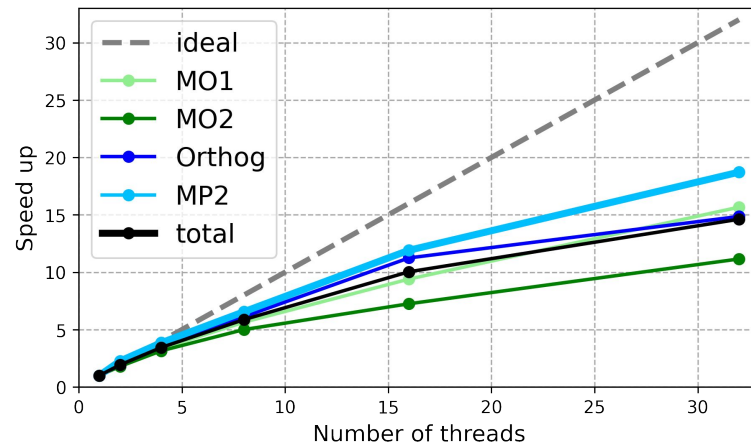

# Performance Benchmarks: OpenMP implementation / strong scaling

#pragma omp parallel for collapse(3)

- reshaping

#pragma omp parallel for collapse(2) reduction(+:E<sub>MP2</sub>)

- DGEMM + element-wise calculations





# Distributed Memory Parallelization with MPI

- Rank three tensor can be distributed to each rank and the change of basis can be done separately on each node.

$$(pq|P) = \sum_q^{N+M} (C^\dagger)_{q\nu} \sum_p^{N+M} (C^\dagger)_{p\mu} (\mu\nu|P)$$

- Orthogonalization can also be done separately on each node after some data transfer

$$B_{pq}^P = \sum_P^{NAUX} (pq|Q) (P|Q)^{-1/2}$$





# Distributed Memory Parallelization with MPI

---



## Algorithm 4: Non-blocking MPI Implementation

---

```
Initialize 1D Cartesian Grid;  
for  $k = 1; k < NumRanks$  do  
  MPI_Isend(B) to the RIGHT neighbor;  
  MPI_Irecv(Brecv_) from the LEFT neighbor ;  
  MP2 Energy Calculation kernel (3) + Update index ;  
  MPI_Waitall(2);  
  Swap B and Brecv_ ;  
end  
MP2 Energy Calculation kernel (3);  
MPI_Reduce(&MP2_partial, &MP2_energy, root rank);
```

---

- Allocate extra memory buffer to do compute/transfer overlap
- The number of MPI Communication is minimized.
- Need an extra bookkeeping index to get the correct MO energy.





# Weak scaling analysis for MPI Implementation

Number of processors	1	2	4	8	16	32
Time for small test	0.17	0.34	0.68	1.50	3.07	6.17
Time for large test	5.37	10.56	23.17	47.54	-	-

- Each processor is map to one core on one node
- Runtime needs to be normalized for weak scaling analysis

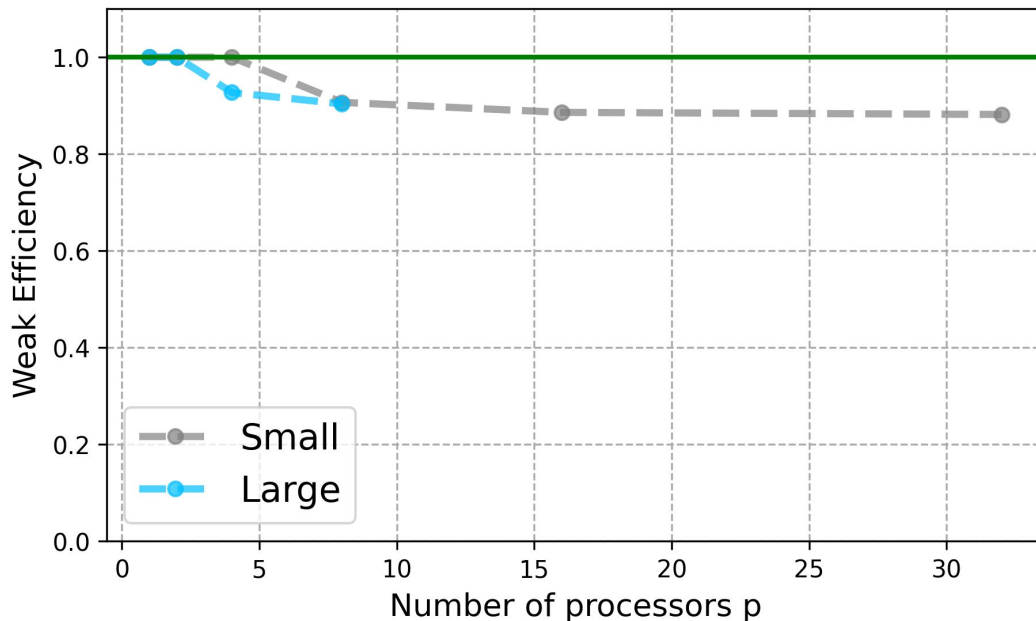




# Weak scaling analysis for MPI Implementation

Number of processors	1	2	4	8	16	32
Time for small test	0.17	0.34	0.68	1.50	3.07	6.17
Time for large test	5.37	10.56	23.17	47.54	-	-

- Each processor is map to one core on one node
- Runtime needs to be normalized for weak scaling analysis





# Final thoughts

- **Integrate BLAS/OpenMP Implementation** into MPI implementation
- Modify the MPI implementation to **divide the work between different processors.**
- Continue working on optimizing the vectorized kernel for MP2 energy calculation.



# Final thoughts

- **Integrate BLAS/OpenMP Implementation** into MPI implementation
  - Modify the MPI implementation to **divide the work between different processors.**
  - Continue working on optimizing the vectorized kernel for MP2 energy calculation.
- 
- Ab initio calculation requires a lot **computing resources**
  - MPI implementation is useful when we do ab initio calculation for **solid state system** (discretize k point in the 1st Brillouin zone)